

02_Tensors_in_PyTorch

June 26, 2025

```
[3]: import torch
      print(torch.__version__)
      print(torch.cuda.is_available())
```

2.6.0+cu124
False

```
[5]: if torch.cuda.is_available():
      device = torch.device("GPU")
      else:
          device = torch.device("cpu")
      print(device)
```

cpu

1 Creating Tensor

```
[9]: # using empty
      a = torch.empty(2,2)
      a
```

```
[9]: tensor([[4.8144e+17, 4.5399e-41],
            [4.8144e+17, 4.5399e-41]])
```

```
[11]: # check type
       type(a)
```

```
[11]: torch.Tensor
```

```
[12]: # using zeros
       torch.zeros(2,3)
```

```
[12]: tensor([[0., 0., 0.],
            [0., 0., 0.]])
```

```
[13]: # using ones
       torch.ones(2,4)
```

```
[13]: tensor([[1., 1., 1., 1.],
            [1., 1., 1., 1.]])
```

```
[16]: # using rand
      torch.rand(2,3)
```

```
[16]: tensor([[0.0886, 0.5391, 0.1133],
            [0.0390, 0.8530, 0.0826]])
```

```
[21]: # use of seed
      torch.manual_seed(100)
      torch.rand(2,3)
```

```
[21]: tensor([[0.1117, 0.8158, 0.2626],
            [0.4839, 0.6765, 0.7539]])
```

```
[23]: # using tensor
      torch.tensor([[1,2,3],[4,5,6]])
```

```
[23]: tensor([[1, 2, 3],
            [4, 5, 6]])
```

```
[40]: # other ways
      # arange
      print(f'using arange ->{torch.arange(0,10,2)}')
      print('-----')
      # linspace
      print(f'using linspace ->{torch.linspace(1,10,10)}')
      print('-----')
      # eye
      print(f'using eye ->{torch.eye(7)}')
      print('-----')
      # diag
      print(f'using diag ->{torch.diag(torch.rand(3,4))}')
      print('-----')
      # full
      print(f'using full ->{torch.full((2,3),6)}')
```

```
using arange ->tensor([0, 2, 4, 6, 8])
```

```
-----
```

```
using linspace ->tensor([ 1.,  2.,  3.,  4.,  5.,  6.,  7.,  8.,  9., 10.]
```

```
-----
```

```
using eye ->tensor([[1., 0., 0., 0., 0., 0., 0.],
                  [0., 1., 0., 0., 0., 0., 0.],
                  [0., 0., 1., 0., 0., 0., 0.],
                  [0., 0., 0., 1., 0., 0., 0.],
                  [0., 0., 0., 0., 1., 0., 0.],
                  [0., 0., 0., 0., 0., 1., 0.]])
```

```

[0., 0., 0., 0., 0., 0., 1.])
-----
using diag ->tensor([0.3039, 0.7590, 0.9454])
-----
using full ->tensor([[6, 6, 6],
                    [6, 6, 6]])

```

2 Tensor Shapes

```
[54]: x = torch.tensor([[1,3,5],[2,4,6]])
      x
```

```
[54]: tensor([[1, 3, 5],
              [2, 4, 6]])
```

```
[45]: x.shape
```

```
[45]: torch.Size([2, 3])
```

```
[49]: torch.empty_like(x)
```

```
[49]: tensor([[ 139149907971472,  139149907971472,          0],
              [          0,          0, 7310593858020254331]])
```

```
[50]: torch.zeros_like(x)
```

```
[50]: tensor([[0, 0, 0],
              [0, 0, 0]])
```

```
[51]: torch.ones_like(x)
```

```
[51]: tensor([[1, 1, 1],
              [1, 1, 1]])
```

```
[66]: torch.rand_like(x, dtype=torch.float32)
```

```
[66]: tensor([[0.1893, 0.9186, 0.2131],
              [0.3957, 0.6017, 0.4234]])
```

3 Tensor Data Types

```
[57]: # find data types
      x.dtype
```

```
[57]: torch.int64
```

```
[59]: # assign data type
      torch.tensor([1,2,3],dtype=torch.float64)
```

```
[59]: tensor([1., 2., 3.], dtype=torch.float64)
```

```
[60]: torch.tensor([1.4,4.2,5.3],dtype=torch.int64)
```

```
[60]: tensor([1, 4, 5])
```

```
[61]: # using to()
      x.to(torch.float64)
```

```
[61]: tensor([[1., 3., 5.],
            [2., 4., 6.]], dtype=torch.float64)
```

Data Type	Dtype	Description
32-bit Floating Point	<code>torch.float32</code>	Standard floating-point type used for most deep learning tasks. Provides a balance between precision and memory usage.
64-bit Floating Point	<code>torch.float64</code>	Double-precision floating point. Useful for high-precision numerical tasks but uses more memory.
16-bit Floating Point	<code>torch.float16</code>	Half-precision floating point. Commonly used in mixed-precision training to reduce memory and computational overhead on modern GPUs.
BFloat16	<code>torch.bfloat16</code>	Bfloat16 floating-point format with reduced precision compared to <code>float16</code> . Used in mixed-precision training, especially on TPUs.
8-bit Floating Point	<code>torch.float8_e4m3fn</code>	Ultra-low-precision floating point. Used for experimental applications and extreme memory-constrained environments (less common).
8-bit Integer	<code>torch.int8</code>	8-bit signed integer. Used for quantized models to save memory and computation in inference.
16-bit Integer	<code>torch.int16</code>	16-bit signed integer. Useful for special numerical tasks requiring intermediate precision.
32-bit Integer	<code>torch.int32</code>	Standard signed integer type. Commonly used for indexing and general-purpose numerical tasks.
64-bit Integer	<code>torch.int64</code>	Long integer type. Often used for large indexing arrays or for tasks involving large numbers.
8-bit Unsigned Integer	<code>torch.uint8</code>	8-bit unsigned integer. Commonly used for image data (e.g., pixel values between 0 and 255).
Boolean	<code>torch.bool</code>	Boolean type, stores <code>True</code> or <code>False</code> values. Often used for masks in logical operations.
Complex 64	<code>torch.complex64</code>	Complex number type with 32-bit real and 32-bit imaginary parts. Used for scientific and signal processing tasks.

Data Type	Dtype	Description
Complex 128	<code>torch.complex128</code>	Complex number type with 64-bit real and 64-bit imaginary parts. Offers higher precision but uses more memory.
Quantized Integer	<code>torch.qint8</code>	Quantized signed 8-bit integer. Used in quantized models for efficient inference.
Quantized Unsigned Integer	<code>torch.quint8</code>	Quantized unsigned 8-bit integer. Often used for quantized tensors in image-related tasks.

4 Mathematical operations

4.1 1. Scalar operation

```
[69]: x = torch.rand(2,2)
      x
```

```
[69]: tensor([[0.3079, 0.6269],
            [0.8277, 0.6594]])
```

```
[74]: # addition
      x + 2

      # subtraction
      x - 2

      # multiplication
      x * 3

      # division
      (x * 100)//3

      # mod
      ((x*100)//3)%2
```

```
[74]: tensor([[0., 0.],
            [1., 1.]])
```

```
[75]: # power
      x**2
```

```
[75]: tensor([[0.0948, 0.3931],
            [0.6850, 0.4347]])
```

4.2 2. Element wise operations

```
[78]: a = torch.rand(2,3)
      b = torch.rand(2,3)

      print(a)
      print(b)
```

```
tensor([[0.4847, 0.9436, 0.3904],
        [0.2499, 0.3206, 0.9753]])
tensor([[0.7582, 0.6688, 0.2651],
        [0.2336, 0.5057, 0.5688]])
```

```
[79]: # add
      a + b
      # sub
      a - b
      # multi
      a * b
      # div
      a / b
      # power
      a ** b
      # mod
      a % b
```

```
[79]: tensor([[0.4847, 0.2748, 0.1253],
              [0.0163, 0.3206, 0.4064]])
```

```
[81]: c = torch.tensor([1, -2, 3, -4])
      c
```

```
[81]: tensor([ 1, -2,  3, -4])
```

```
[82]: # absolute
      torch.abs(c)
```

```
[82]: tensor([1, 2, 3, 4])
```

```
[84]: # neg
      torch.neg(c)
```

```
[84]: tensor([-1,  2, -3,  4])
```

```
[86]: d = torch.tensor([1.9,3.5,9.3,2.6,8.4])
      d
```

```
[86]: tensor([1.9000, 3.5000, 9.3000, 2.6000, 8.4000])
```

```
[87]: # round
torch.round(d)
```

```
[87]: tensor([2., 4., 9., 3., 8.])
```

```
[90]: # ceil
torch.ceil(d)
```

```
[90]: tensor([ 2.,  4., 10.,  3.,  9.])
```

```
[91]: # floor
torch.floor(d)
```

```
[91]: tensor([1., 3., 9., 2., 8.])
```

```
[92]: # clamp
torch.clamp(d, min=3, max=8)
```

```
[92]: tensor([3.0000, 3.5000, 8.0000, 3.0000, 8.0000])
```

5 Reduction operation

```
[95]: e = torch.randint(size=(2,3), low=0, high=10)
e
```

```
[95]: tensor([[2, 2, 3],
          [7, 1, 5]])
```

```
[97]: # sum
torch.sum(e)
```

```
[97]: tensor(20)
```

```
[98]: # sum along columns
torch.sum(e, dim=0)
```

```
[98]: tensor([9, 3, 8])
```

```
[99]: # sum along rows
torch.sum(e, dim=1)
```

```
[99]: tensor([ 7, 13])
```

```
[100]: # mean
torch.mean(e.float(), dim=1)
```

```
[100]: tensor([2.3333, 4.3333])
```

```
[101]: # mean along columns  
torch.mean(e.float(), dim=0)
```

```
[101]: tensor([4.5000, 1.5000, 4.0000])
```

```
[102]: # median  
torch.median(e.float())
```

```
[102]: tensor(2.)
```

```
[103]: # max and min  
torch.max(e)  
torch.min(e)
```

```
[103]: tensor(1)
```

```
[104]: # product  
torch.prod(e)
```

```
[104]: tensor(420)
```

```
[106]: # Standard devaiiton  
torch.std(e.float())
```

```
[106]: tensor(2.2509)
```

```
[107]: # variance  
torch.var(e.float())
```

```
[107]: tensor(5.0667)
```

```
[108]: # argmax  
torch.argmax(e)
```

```
[108]: tensor(3)
```

```
[109]: # argmin  
torch.argmin(e)
```

```
[109]: tensor(4)
```

6 Matrix Operations


```
[112]: f = torch.randint(size=(2,3), low=0, high=10)
      g = torch.randint(size=(3,2), low=0, high=10)
```

```
print(f)
print(g)
```

```
tensor([[3, 6, 7],
        [8, 1, 2]])
tensor([[3, 2],
        [7, 9],
        [3, 9]])
```

```
[115]: # matrix multiplication
      torch.matmul(f,g)
```

```
[115]: tensor([[ 72, 123],
              [ 37,  43]])
```

```
[116]: vector1 = torch.tensor([1,2,3])
      vector2 = torch.tensor([4,5,6])
```

```
print(vector1)
print(vector2)
```

```
tensor([1, 2, 3])
tensor([4, 5, 6])
```

```
[117]: # dot product
      torch.dot(vector1, vector2)
```

```
[117]: tensor(32)
```

```
[119]: # transpose
      print(f)
      torch.t(f)
```

```
tensor([[3, 6, 7],
        [8, 1, 2]])
```

```
[119]: tensor([[3, 8],
              [6, 1],
              [7, 2]])
```

```
[127]: h = torch.randint(size=(3,3), low=0, high=10, dtype=torch.float32)
      h
```

```
[127]: tensor([[5., 0., 6.],
              [7., 7., 8.],
              [8., 9., 6.]])
```

```
[128]: # determinant
torch.det(h)
```

```
[128]: tensor(-108.0000)
```

```
[129]: # inverse
torch.inverse(h)
```

```
[129]: tensor([[ 0.2778, -0.5000,  0.3889],
          [-0.2037,  0.1667, -0.0185],
          [-0.0648,  0.4167, -0.3241]])
```

6.1 Comparison operations

```
[130]: i = torch.randint(size=(2,3), low=0, high=10)
j = torch.randint(size=(2,3), low=0, high=10)

print(i)
print(j)
```

```
tensor([[2, 9, 4],
        [5, 5, 3]])
tensor([[5, 0, 7],
        [4, 7, 1]])
```

```
[131]: # greater than
i > j
```

```
[131]: tensor([[False,  True, False],
          [ True, False,  True]])
```

```
[132]: # less than
i < j
```

```
[132]: tensor([[ True, False,  True],
          [False,  True, False]])
```

```
[133]: # equal to
i == j
```

```
[133]: tensor([[False, False, False],
          [False, False, False]])
```

```
[134]: # not equal to
i != j
```

```
[134]: tensor([[ True,  True,  True],
          [ True,  True,  True]])
```

```
[137]: # greater than equal to  
i >= j
```

```
[137]: tensor([[False,  True, False],  
          [ True, False,  True]])
```

```
[138]: # less than equal to  
i <= j
```

```
[138]: tensor([[ True, False,  True],  
          [False,  True, False]])
```

6.2 6. Special functions

```
[146]: k = torch.randint(size=(2,3), low=0, high=10, dtype=torch.float32)  
k
```

```
[146]: tensor([[0., 8., 5.],  
          [1., 3., 8.]])
```

```
[140]: # log  
torch.log(k)
```

```
[140]: tensor([[1.0986, 1.7918, 2.1972],  
          [1.9459,  -inf, 0.0000]])
```

```
[141]: # exp  
torch.exp(k)
```

```
[141]: tensor([[2.0086e+01, 4.0343e+02, 8.1031e+03],  
          [1.0966e+03, 1.0000e+00, 2.7183e+00]])
```

```
[142]: # sqrt  
torch.sqrt(k)
```

```
[142]: tensor([[1.7321, 2.4495, 3.0000],  
          [2.6458, 0.0000, 1.0000]])
```

```
[143]: # sigmoid  
torch.sigmoid(k)
```

```
[143]: tensor([[0.9526, 0.9975, 0.9999],  
          [0.9991, 0.5000, 0.7311]])
```

```
[147]: # softmax  
torch.softmax(k, dim=0)
```

```
[147]: tensor([[0.2689, 0.9933, 0.0474],
              [0.7311, 0.0067, 0.9526]])
```

```
[148]: # relu
       torch.relu(k)
```

```
[148]: tensor([[0., 8., 5.],
              [1., 3., 8.]])
```

7 Inplace Operations

```
[149]: m = torch.rand(2,3)
       n = torch.rand(2,3)

       print(m)
       print(n)

       tensor([[0.9355, 0.1430, 0.3933],
              [0.1124, 0.3087, 0.9973]])
       tensor([[0.4257, 0.6890, 0.9657],
              [0.0257, 0.4205, 0.0656]])
```

```
[150]: m + n
```

```
[150]: tensor([[1.3612, 0.8320, 1.3590],
              [0.1382, 0.7292, 1.0629]])
```

```
[151]: m.add_(n)
```

```
[151]: tensor([[1.3612, 0.8320, 1.3590],
              [0.1382, 0.7292, 1.0629]])
```

```
[152]: m
```

```
[152]: tensor([[1.3612, 0.8320, 1.3590],
              [0.1382, 0.7292, 1.0629]])
```

```
[153]: n
```

```
[153]: tensor([[0.4257, 0.6890, 0.9657],
              [0.0257, 0.4205, 0.0656]])
```

```
[154]: torch.relu(m)
```

```
[154]: tensor([[1.3612, 0.8320, 1.3590],
              [0.1382, 0.7292, 1.0629]])
```

```
[158]: m.relu_()
```

```
[158]: tensor([[1.3612, 0.8320, 1.3590],
              [0.1382, 0.7292, 1.0629]])
```

```
[157]: m
```

```
[157]: tensor([[1.3612, 0.8320, 1.3590],
              [0.1382, 0.7292, 1.0629]])
```

8 Copying a tensor

```
[159]: a = torch.rand(2,3)
a
```

```
[159]: tensor([[0.4508, 0.0553, 0.3140],
              [0.7460, 0.9357, 0.8925]])
```

```
[160]: b = a
```

```
[161]: b
```

```
[161]: tensor([[0.4508, 0.0553, 0.3140],
              [0.7460, 0.9357, 0.8925]])
```

```
[164]: a[0][0] = 0
a, b
```

```
[164]: (tensor([[0.0000, 0.0553, 0.3140],
              [0.7460, 0.9357, 0.8925]]),
       tensor([[0.0000, 0.0553, 0.3140],
              [0.7460, 0.9357, 0.8925]]))
```

```
[166]: id(a) , id(b)
```

```
[166]: (139145171801808, 139145171801808)
```

```
[168]: b = a.clone()
a,b
```

```
[168]: (tensor([[0.0000, 0.0553, 0.3140],
              [0.7460, 0.9357, 0.8925]]),
       tensor([[0.0000, 0.0553, 0.3140],
              [0.7460, 0.9357, 0.8925]]))
```

```
[169]: id(a), id(b)
```

```
[169]: (139145171801808, 139145170351120)
```

```
[171]: a[0][0] = 9  
a , b
```

```
[171]: (tensor([[9.0000, 0.0553, 0.3140],  
               [0.7460, 0.9357, 0.8925]]),  
       tensor([[0.0000, 0.0553, 0.3140],  
               [0.7460, 0.9357, 0.8925]]))
```