

10_With_Hyperparameteres_ANN_on_GPU_fashion_mnist_PyTorch

July 14, 2025

1 ANN with PyTorch - Fashion MNIST

2 Import libraries

```
[32]: # import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from sklearn.model_selection import train_test_split
```

```
[30]: # Set random seed for reproducibility
torch.manual_seed(42)
```

```
[30]: <torch._C.Generator at 0x7c096adf0f30>
```

```
[31]: # Check for GPU
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"Using device: {device}")
```

Using device: cuda

3 import dataset

```
[7]: !unzip /content/fashion_mnist_dataset.zip
```

```
Archive: /content/fashion_mnist_dataset.zip
  inflating: fashion-mnist_test.csv
  inflating: fashion-mnist_train.csv
  inflating: t10k-images-idx3-ubyte
  inflating: t10k-labels-idx1-ubyte
  inflating: train-images-idx3-ubyte
  inflating: train-labels-idx1-ubyte
```

```
[33]: # import dataset
df = pd.read_csv('/content/fashion-mnist_train.csv')
df.head()
```

```
[33]:
```

| | label | pixel1 | pixel2 | pixel3 | pixel4 | pixel5 | pixel6 | pixel7 | pixel8 | \ |
|---|-------|--------|--------|--------|--------|--------|--------|--------|--------|---|
| 0 | 2 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 1 | 9 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 2 | 6 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 5 | |
| 3 | 0 | 0 | 0 | 0 | 1 | 2 | 0 | 0 | 0 | |
| 4 | 3 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |


| | pixel19 | ... | pixel775 | pixel776 | pixel777 | pixel778 | pixel779 | pixel780 | \ |
|---|---------|-----|----------|----------|----------|----------|----------|----------|---|
| 0 | 0 | ... | 0 | 0 | 0 | 0 | 0 | 0 | |
| 1 | 0 | ... | 0 | 0 | 0 | 0 | 0 | 0 | |
| 2 | 0 | ... | 0 | 0 | 0 | 30 | 43 | 0 | |
| 3 | 0 | ... | 3 | 0 | 0 | 0 | 0 | 1 | |
| 4 | 0 | ... | 0 | 0 | 0 | 0 | 0 | 0 | |

| | pixel781 | pixel782 | pixel783 | pixel784 |
|---|----------|----------|----------|----------|
| 0 | 0 | 0 | 0 | 0 |
| 1 | 0 | 0 | 0 | 0 |
| 2 | 0 | 0 | 0 | 0 |
| 3 | 0 | 0 | 0 | 0 |
| 4 | 0 | 0 | 0 | 0 |

[5 rows x 785 columns]

```
[34]: # let's plot first 10 images with 5x5 grid
fig, ax = plt.subplots(2, 5, figsize=(15, 10))
for i in range(10):
    ax[i//5, i%5].imshow(df.iloc[i, 1:].values.reshape(28, 28), cmap='gray')
    ax[i//5, i%5].axis('off')
    ax[i//5, i%5].set_title(df.iloc[i, 0])

plt.tight_layout()
plt.show()
```



3.1 Split the data

```
[35]: # split the data into features and labels
X = df.iloc[:, 1:].values
y = df.iloc[:, 0].values
X.shape, y.shape
```

```
[35]: ((60000, 784), (60000,))
```

```
[36]: # split the data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
[36]: ((48000, 784), (12000, 784), (48000,), (12000,))
```

3.2 Scale the data

```
[37]: # scaling the data
X_train = X_train / 255.0
X_test = X_test / 255.0
```

4 create dataloader

```
[38]: # create CustomDataset class
class CustomDataset(torch.utils.data.Dataset):
    def __init__(self, X, y):
        self.X = torch.tensor(X, dtype=torch.float32)
```

```

        self.y = torch.tensor(y, dtype=torch.long)

    def __len__(self):
        return len(self.X)

    def __getitem__(self, index):
        x = self.X[index]
        y = self.y[index]
        return x, y

```

```

[39]: # create dataloader
train_dataset = CustomDataset(X_train, y_train)
test_dataset = CustomDataset(X_test, y_test)

```

5 Model building

```

[47]: class MyNN(nn.Module):

    def __init__(self, input_dim, output_dim, num_hidden_layers,
↳neurons_per_layer, dropout_rate):
        super().__init__()
        layers = []
        for i in range(num_hidden_layers):

            layers.append(nn.Linear(input_dim, neurons_per_layer))
            layers.append(nn.BatchNorm1d(neurons_per_layer))
            layers.append(nn.ReLU())
            layers.append(nn.Dropout(dropout_rate))
            input_dim = neurons_per_layer

        layers.append(nn.Linear(neurons_per_layer, output_dim))

        self.model = nn.Sequential(*layers)

    def forward(self, x):
        x = self.model(x)
        return x

```

6 Objective function

```

[51]: # Objective function
def objective(trial):
    # next hyperparameter values from the search space
    num_hidden_layers = trial.suggest_int('num_hidden_layers', 1, 5)
    neurons_per_layer = trial.suggest_int('neurons_per_layer', 8, 128, step=8)

```

```

epochs = trial.suggest_int('epochs', 10, 100, step=5)
lr = trial.suggest_float('lr', 1e-5, 1e-1, log=True)
dropout_rate = trial.suggest_float('dropout_rate', 0.1, 0.5, step=0.1)
batch_size = trial.suggest_categorical('batch_size', [32, 64, 128])
optimizer_name = trial.suggest_categorical('optimizer', ['Adam', 'SGD',
↪'RMSprop'])
weight_decay = trial.suggest_float('weight_decay', 1e-5, 1e-3, log=True)

# data loader
train_loader = torch.utils.data.DataLoader(train_dataset,
↪batch_size=batch_size, shuffle=True, pin_memory=True)
test_loader = torch.utils.data.DataLoader(test_dataset,
↪batch_size=batch_size, shuffle=False, pin_memory=True)

# model init
input_dim = X_train.shape[1]
output_dim = 10
model = MyNN(input_dim, output_dim, num_hidden_layers, neurons_per_layer,
↪dropout_rate)
model.to(device)

# optimizer selection
criterion = nn.CrossEntropyLoss()

if optimizer_name == 'Adam':
    optimizer = optim.Adam(model.parameters(), lr=lr, weight_decay=weight_decay)

elif optimizer_name == 'SGD':
    optimizer = optim.SGD(model.parameters(), lr=lr, momentum=0.9,
↪weight_decay=weight_decay)

elif optimizer_name == 'RMSprop':
    optimizer = optim.RMSprop(model.parameters(), lr=lr,
↪weight_decay=weight_decay)

# training loop
for epoch in range(epochs):

    for batch_features, batch_labels in train_loader:
        # move data to gpu
        batch_features = batch_features.to(device)
        batch_labels = batch_labels.to(device)

        # forward pass
        outputs = model(batch_features)

```

```

        # calculate loss
        loss = criterion(outputs, batch_labels)

        # zero gradients
        optimizer.zero_grad()

        # backward pass
        loss.backward()

        # update weights
        optimizer.step()

    # evaluation
    model.eval()

    with torch.no_grad():
        correct = 0
        total = 0
        for images, labels in test_loader:
            images = images.to(device)
            labels = labels.to(device)
            outputs = model(images)
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

        accuracy = correct / total

    return accuracy

```

```

[52]: # lets create study
import optuna
study = optuna.create_study(direction='maximize')

```

[I 2025-07-08 08:33:25,359] A new study created in memory with name: no-name-d67083ee-7fe1-408f-945e-21864b5f1b97

```

[53]: study.optimize(objective, n_trials=20)

```

[I 2025-07-08 08:36:05,974] Trial 0 finished with value: 0.8814166666666666 and parameters: {'num_hidden_layers': 3, 'neurons_per_layer': 72, 'epochs': 70, 'lr': 0.0016976313223769004, 'dropout_rate': 0.4, 'batch_size': 64, 'optimizer': 'Adam', 'weight_decay': 5.4260526198393654e-05}. Best is trial 0 with value: 0.8814166666666666.

[I 2025-07-08 08:37:10,028] Trial 1 finished with value: 0.87925 and parameters: {'num_hidden_layers': 5, 'neurons_per_layer': 64, 'epochs': 40, 'lr': 0.001687930783405996, 'dropout_rate': 0.2, 'batch_size': 128, 'optimizer': 'Adam', 'weight_decay': 0.00014998256990705265}. Best is trial 0 with value: 0.8814166666666666.

0.8814166666666666.

[I 2025-07-08 08:42:56,455] Trial 2 finished with value: 0.8005833333333333 and parameters: {'num_hidden_layers': 2, 'neurons_per_layer': 112, 'epochs': 100, 'lr': 0.007960225498043969, 'dropout_rate': 0.5, 'batch_size': 32, 'optimizer': 'RMSprop', 'weight_decay': 0.00046980805192913753}. Best is trial 0 with value: 0.8814166666666666.

[I 2025-07-08 08:45:05,165] Trial 3 finished with value: 0.86425 and parameters: {'num_hidden_layers': 5, 'neurons_per_layer': 72, 'epochs': 85, 'lr': 0.004344539222441111, 'dropout_rate': 0.5, 'batch_size': 128, 'optimizer': 'SGD', 'weight_decay': 0.0005744733676755499}. Best is trial 0 with value: 0.8814166666666666.

[I 2025-07-08 08:46:17,479] Trial 4 finished with value: 0.8590833333333333 and parameters: {'num_hidden_layers': 3, 'neurons_per_layer': 64, 'epochs': 60, 'lr': 0.0009264008464073184, 'dropout_rate': 0.5, 'batch_size': 128, 'optimizer': 'SGD', 'weight_decay': 3.8621196692052926e-05}. Best is trial 0 with value: 0.8814166666666666.

[I 2025-07-08 08:47:42,236] Trial 5 finished with value: 0.8810833333333333 and parameters: {'num_hidden_layers': 2, 'neurons_per_layer': 120, 'epochs': 80, 'lr': 0.0010222542982193145, 'dropout_rate': 0.5, 'batch_size': 128, 'optimizer': 'SGD', 'weight_decay': 0.00018070933916449335}. Best is trial 0 with value: 0.8814166666666666.

[I 2025-07-08 08:52:01,570] Trial 6 finished with value: 0.8765 and parameters: {'num_hidden_layers': 1, 'neurons_per_layer': 128, 'epochs': 100, 'lr': 0.009017944278052047, 'dropout_rate': 0.1, 'batch_size': 32, 'optimizer': 'SGD', 'weight_decay': 0.00020421724797268796}. Best is trial 0 with value: 0.8814166666666666.

[I 2025-07-08 08:54:35,361] Trial 7 finished with value: 0.85775 and parameters: {'num_hidden_layers': 1, 'neurons_per_layer': 32, 'epochs': 60, 'lr': 0.008589890737706248, 'dropout_rate': 0.4, 'batch_size': 32, 'optimizer': 'SGD', 'weight_decay': 0.00030725793580119805}. Best is trial 0 with value: 0.8814166666666666.

[I 2025-07-08 08:56:51,854] Trial 8 finished with value: 0.8886666666666667 and parameters: {'num_hidden_layers': 5, 'neurons_per_layer': 80, 'epochs': 90, 'lr': 0.006949968609931943, 'dropout_rate': 0.30000000000000004, 'batch_size': 128, 'optimizer': 'SGD', 'weight_decay': 6.677334308862196e-05}. Best is trial 8 with value: 0.8886666666666667.

[I 2025-07-08 08:57:05,253] Trial 9 finished with value: 0.8750833333333333 and parameters: {'num_hidden_layers': 1, 'neurons_per_layer': 48, 'epochs': 15, 'lr': 0.014994162103390576, 'dropout_rate': 0.2, 'batch_size': 128, 'optimizer': 'SGD', 'weight_decay': 6.916027568905143e-05}. Best is trial 8 with value: 0.8886666666666667.

```
[54]: study.best_value, study.best_params
```

```
[54]: (0.8886666666666667,
      {'num_hidden_layers': 5,
       'neurons_per_layer': 80,
```

```

'epochs': 90,
'lr': 0.006949968609931943,
'dropout_rate': 0.30000000000000004,
'batch_size': 128,
'optimizer': 'SGD',
'weight_decay': 6.677334308862196e-05})

```

```

[56]: # Best trial se model banaen
best_trial = study.best_trial
best_params = best_trial.params
# model init
input_dim = X_train.shape[1]
output_dim = 10
# Model ko train karein best hyperparameters ke saath
model = MyNN(input_dim, output_dim, best_params['num_hidden_layers'],
             ↪best_params['neurons_per_layer'], best_params['dropout_rate'])
# ... (training code jaisa pehle diya)

# Model ko save karein
torch.save(model.state_dict(), 'best_model.pth')

```